

# HandWritten Digit Recognition

First DeepLearning project

In this project, we are using mnist dataset and try to train a neural network to do the digit classification

```
In [33]: import numpy as np
import pandas as pd
import matplotlib as plt
import tensorflow as tf
from tensorflow import keras
import matplotlib.pyplot as plt

print(tf.__version__)
```

2.13.0

load the mnist data

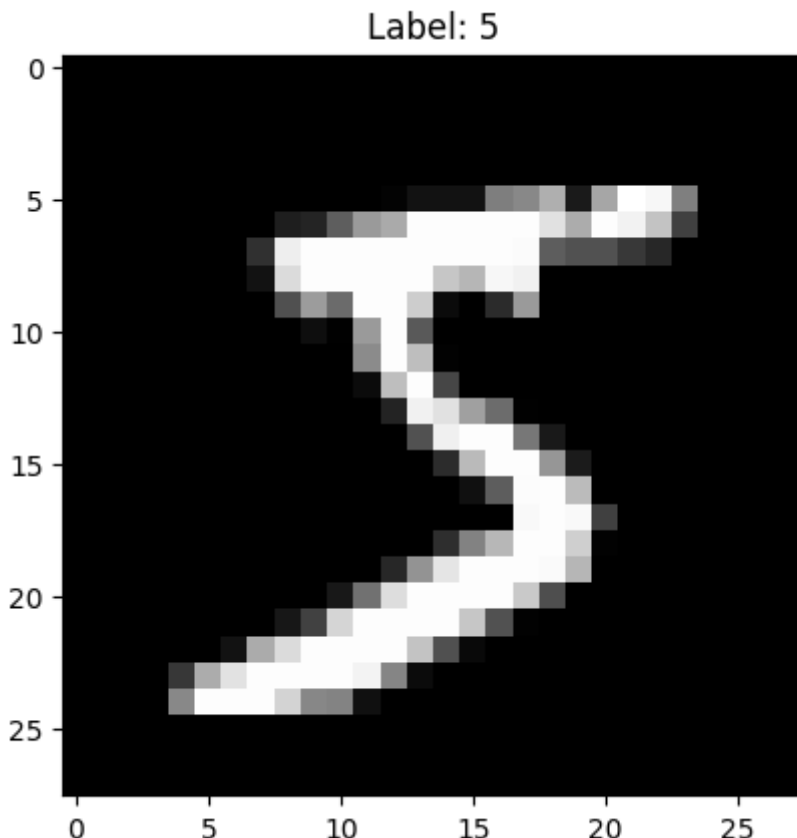
```
In [2]: from tensorflow.keras.datasets import mnist
```

## load the mnist dataset

```
In [3]: (train_images, train_labels), (test_images, test_labels) = mnist.load_data()
```

## Inspect the data

```
In [4]: plt.imshow(train_images[0], cmap='gray')  
plt.title(f"Label: {train_labels[0]}")  
plt.show()
```



before we do anything, we have to remember to normalize the data

## Normalize the images to be values between 0 and 1

```
In [5]: train_images = train_images.astype('float32') / 255  
test_images = test_images.astype('float32') / 255
```

then,

## One-hot encode the labels

```
In [6]: train_labels = tf.keras.utils.to_categorical(train_labels)  
test_labels = tf.keras.utils.to_categorical(test_labels)
```

to understand how the data is structured, we print it

```
In [7]: print(train_images)
```

```
[[[0. 0. 0. ... 0. 0. 0.]
  [0. 0. 0. ... 0. 0. 0.]
  [0. 0. 0. ... 0. 0. 0.]
  ...
  [0. 0. 0. ... 0. 0. 0.]
  [0. 0. 0. ... 0. 0. 0.]
  [0. 0. 0. ... 0. 0. 0.]]

[[[0. 0. 0. ... 0. 0. 0.]
  [0. 0. 0. ... 0. 0. 0.]
  [0. 0. 0. ... 0. 0. 0.]
  ...
  [0. 0. 0. ... 0. 0. 0.]
  [0. 0. 0. ... 0. 0. 0.]
  [0. 0. 0. ... 0. 0. 0.]]

[[[0. 0. 0. ... 0. 0. 0.]
  [0. 0. 0. ... 0. 0. 0.]
  [0. 0. 0. ... 0. 0. 0.]
  ...
  [0. 0. 0. ... 0. 0. 0.]
  [0. 0. 0. ... 0. 0. 0.]
  [0. 0. 0. ... 0. 0. 0.]]

...

[[[0. 0. 0. ... 0. 0. 0.]
  [0. 0. 0. ... 0. 0. 0.]
  [0. 0. 0. ... 0. 0. 0.]
  ...
  [0. 0. 0. ... 0. 0. 0.]
  [0. 0. 0. ... 0. 0. 0.]
  [0. 0. 0. ... 0. 0. 0.]]

[[[0. 0. 0. ... 0. 0. 0.]
  [0. 0. 0. ... 0. 0. 0.]
  [0. 0. 0. ... 0. 0. 0.]
  ...
  [0. 0. 0. ... 0. 0. 0.]
  [0. 0. 0. ... 0. 0. 0.]
  [0. 0. 0. ... 0. 0. 0.]]

[[[0. 0. 0. ... 0. 0. 0.]
  [0. 0. 0. ... 0. 0. 0.]
  [0. 0. 0. ... 0. 0. 0.]
  ...
  [0. 0. 0. ... 0. 0. 0.]
  [0. 0. 0. ... 0. 0. 0.]
  [0. 0. 0. ... 0. 0. 0.]]]
```

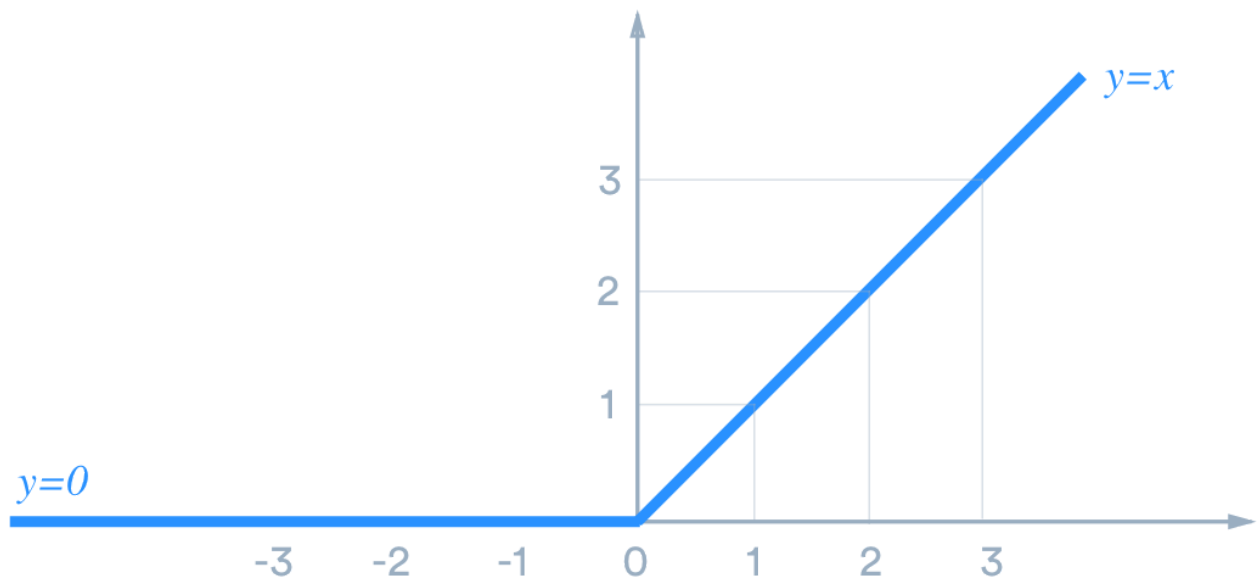
```
In [8]: print(train_labels.shape)
print(train_images.shape)

(60000, 10)
(60000, 28, 28)
```

so each number is 28 pixels x 28 pixels

## 1. Now, we start to set up layers for it

we use relu functions as the activation function since is good and its simple, easy for training and processing data



```
In [9]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Flatten, Dense

model = Sequential([
    Flatten(input_shape=(28, 28)), # Flatten the 28x28 images
    Dense(128, activation='relu'),  # Fully connected layer with 128 units
    Dense(10, activation='softmax') # Output layer with 10 units (for 10 classes)
])
```

## 2. Compile the Model

```
In [10]: model.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])

model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 784)	0
dense (Dense)	(None, 128)	100480
dense_1 (Dense)	(None, 10)	1290
Total params: 101770 (397.54 KB)		
Trainable params: 101770 (397.54 KB)		
Non-trainable params: 0 (0.00 Byte)		

### 3. Train the Model

```
In [11]: history = model.fit(
            train_images,
            train_labels,
            epochs=5,
            batch_size=32,
            validation_data=(test_images, test_labels)
        )
```

```
Epoch 1/5
1875/1875 [=====] - 2s 817us/step - loss: 0.2555
- accuracy: 0.9281 - val_loss: 0.1287 - val_accuracy: 0.9625
Epoch 2/5
1875/1875 [=====] - 1s 761us/step - loss: 0.1095
- accuracy: 0.9682 - val_loss: 0.0963 - val_accuracy: 0.9700
Epoch 3/5
1875/1875 [=====] - 1s 714us/step - loss: 0.0759
- accuracy: 0.9769 - val_loss: 0.0829 - val_accuracy: 0.9726
Epoch 4/5
1875/1875 [=====] - 1s 694us/step - loss: 0.0575
- accuracy: 0.9820 - val_loss: 0.0816 - val_accuracy: 0.9748
Epoch 5/5
1875/1875 [=====] - 1s 726us/step - loss: 0.0447
- accuracy: 0.9858 - val_loss: 0.0749 - val_accuracy: 0.9777
```

### 4. Evaluate the Model

```
In [12]: test_loss, test_acc = model.evaluate(test_images, test_labels)
print(f"Test Accuracy: {test_acc * 100:.2f}%")
```

```
313/313 [=====] - 0s 418us/step - loss: 0.0749 -
accuracy: 0.9777
Test Accuracy: 97.77%
```

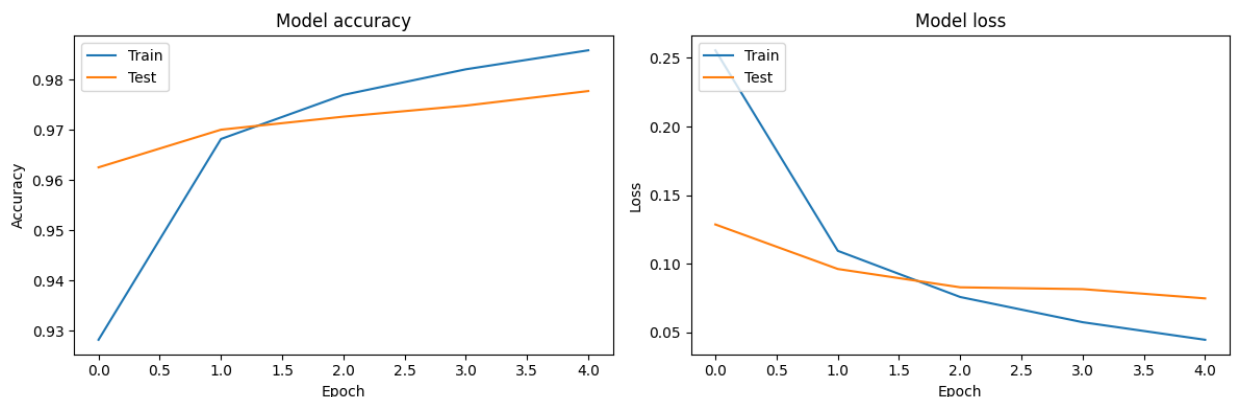
## lets virtualize it

```
In [13]: # Plot training & validation accuracy values
plt.figure(figsize=(12, 4))

plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')

# Plot training & validation loss values
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')

plt.tight_layout()
plt.show()
```



Great! we got 97% accuracy!

we can still do some improvements

```
In [14]: model = Sequential([
    Flatten(input_shape=(28, 28)), # Flatten the 28x28 images
    Dense(256, activation='relu'), # Fully connected layer with 128 units
    Dense(10, activation='softmax') # Output layer with 10 units (for 10 c
])

model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

model.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
flatten_1 (Flatten)	(None, 784)	0
dense_2 (Dense)	(None, 256)	200960
dense_3 (Dense)	(None, 10)	2570

=====  
Total params: 203530 (795.04 KB)  
Trainable params: 203530 (795.04 KB)  
Non-trainable params: 0 (0.00 Byte)

now we change the epochs count to 10, but just in case we add early\_stopping to avoid overfitting

```
In [15]: from tensorflow.keras.callbacks import EarlyStopping
early_stopping = EarlyStopping(
    min_delta=0.001, # minimum amount of change to count as an improvement
    patience=20, # how many epochs to wait before stopping
    restore_best_weights=True,
)

history = model.fit(
    train_images,
    train_labels,
    epochs=10,
    batch_size=32,
    callbacks=[early_stopping], # put your callbacks in a list
    validation_data=(test_images, test_labels)
)
```

Epoch 1/10

1875/1875 [=====] - 2s 837us/step - loss: 0.2265  
- accuracy: 0.9342 - val\_loss: 0.1143 - val\_accuracy: 0.9652

Epoch 2/10

1875/1875 [=====] - 1s 790us/step - loss: 0.0937  
- accuracy: 0.9721 - val\_loss: 0.0896 - val\_accuracy: 0.9692

Epoch 3/10

1875/1875 [=====] - 1s 782us/step - loss: 0.0622  
- accuracy: 0.9815 - val\_loss: 0.0675 - val\_accuracy: 0.9786

Epoch 4/10

1875/1875 [=====] - 1s 795us/step - loss: 0.0438  
- accuracy: 0.9866 - val\_loss: 0.0682 - val\_accuracy: 0.9785

Epoch 5/10

1875/1875 [=====] - 2s 822us/step - loss: 0.0347  
- accuracy: 0.9889 - val\_loss: 0.0706 - val\_accuracy: 0.9780

Epoch 6/10

1875/1875 [=====] - 2s 829us/step - loss: 0.0247  
- accuracy: 0.9922 - val\_loss: 0.0760 - val\_accuracy: 0.9775

Epoch 7/10

1875/1875 [=====] - 2s 801us/step - loss: 0.0196  
- accuracy: 0.9936 - val\_loss: 0.0744 - val\_accuracy: 0.9793

Epoch 8/10

1875/1875 [=====] - 2s 982us/step - loss: 0.0164  
- accuracy: 0.9946 - val\_loss: 0.0748 - val\_accuracy: 0.9802

Epoch 9/10

1875/1875 [=====] - 1s 783us/step - loss: 0.0130  
- accuracy: 0.9956 - val\_loss: 0.0791 - val\_accuracy: 0.9796

Epoch 10/10

1875/1875 [=====] - 1s 779us/step - loss: 0.0109  
- accuracy: 0.9965 - val\_loss: 0.0784 - val\_accuracy: 0.9815

Great! almost 100% accuracy, lets try further with higher epoches



```
In [16]: history = model.fit(
    train_images,
    train_labels,
    epochs=15,
    batch_size=32,
    callbacks=[early_stopping], # put your callbacks in a list
    validation_data=(test_images, test_labels)
)
```

```
Epoch 1/15
1875/1875 [=====] - 2s 838us/step - loss: 0.0098
- accuracy: 0.9968 - val_loss: 0.0725 - val_accuracy: 0.9830
Epoch 2/15
1875/1875 [=====] - 2s 823us/step - loss: 0.0094
- accuracy: 0.9967 - val_loss: 0.0832 - val_accuracy: 0.9801
Epoch 3/15
1875/1875 [=====] - 2s 840us/step - loss: 0.0080
- accuracy: 0.9973 - val_loss: 0.0831 - val_accuracy: 0.9815
Epoch 4/15
1875/1875 [=====] - 2s 836us/step - loss: 0.0070
- accuracy: 0.9976 - val_loss: 0.0913 - val_accuracy: 0.9792
Epoch 5/15
1875/1875 [=====] - 1s 788us/step - loss: 0.0061
- accuracy: 0.9981 - val_loss: 0.0992 - val_accuracy: 0.9785
Epoch 6/15
1875/1875 [=====] - 1s 791us/step - loss: 0.0050
- accuracy: 0.9985 - val_loss: 0.0985 - val_accuracy: 0.9808
Epoch 7/15
1875/1875 [=====] - 1s 785us/step - loss: 0.0072
- accuracy: 0.9976 - val_loss: 0.1078 - val_accuracy: 0.9798
Epoch 8/15
1875/1875 [=====] - 2s 816us/step - loss: 0.0049
- accuracy: 0.9983 - val_loss: 0.1113 - val_accuracy: 0.9797
Epoch 9/15
1875/1875 [=====] - 2s 802us/step - loss: 0.0055
- accuracy: 0.9979 - val_loss: 0.1200 - val_accuracy: 0.9784
Epoch 10/15
1875/1875 [=====] - 1s 787us/step - loss: 0.0059
- accuracy: 0.9979 - val_loss: 0.1132 - val_accuracy: 0.9809
Epoch 11/15
1875/1875 [=====] - 2s 810us/step - loss: 0.0045
- accuracy: 0.9986 - val_loss: 0.1189 - val_accuracy: 0.9804
Epoch 12/15
1875/1875 [=====] - 2s 847us/step - loss: 0.0060
- accuracy: 0.9980 - val_loss: 0.1043 - val_accuracy: 0.9813
Epoch 13/15
1875/1875 [=====] - 2s 819us/step - loss: 0.0022
- accuracy: 0.9993 - val_loss: 0.1362 - val_accuracy: 0.9780
Epoch 14/15
1875/1875 [=====] - 2s 811us/step - loss: 0.0060
- accuracy: 0.9979 - val_loss: 0.1164 - val_accuracy: 0.9798
Epoch 15/15
1875/1875 [=====] - 1s 799us/step - loss: 0.0034
- accuracy: 0.9988 - val_loss: 0.1139 - val_accuracy: 0.9815
```

it seems if I directly use another fit functions, the loss start as where I ended last time, which might suggest that it trains again on the model I have just trained

If we want to avoid it, I guess we should start it over

```
In [28]: model = Sequential([
    Flatten(input_shape=(28, 28)), # Flatten the 28x28 images
    Dense(128, activation='relu'), # Fully connected layer with 128 units
    Dense(10, activation='softmax') # Output layer with 10 units (for 10 c
])

model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
early_stopping = EarlyStopping(
    min_delta=0.001, # minimum amount of change to count as an improvement
    patience=20, # how many epochs to wait before stopping
    restore_best_weights=True,
)
history = model.fit(
    train_images,
    train_labels,
    epochs=10,
    batch_size=32,
    callbacks=[early_stopping], # put your callbacks in a list
    validation_data=(test_images, test_labels)
)
```

Epoch 1/10

1875/1875 [=====] - 1s 696us/step - loss: 0.2561  
- accuracy: 0.9274 - val\_loss: 0.1318 - val\_accuracy: 0.9618

Epoch 2/10

1875/1875 [=====] - 1s 673us/step - loss: 0.1139  
- accuracy: 0.9664 - val\_loss: 0.0943 - val\_accuracy: 0.9717

Epoch 3/10

1875/1875 [=====] - 1s 713us/step - loss: 0.0772  
- accuracy: 0.9766 - val\_loss: 0.0832 - val\_accuracy: 0.9739

Epoch 4/10

1875/1875 [=====] - 1s 660us/step - loss: 0.0570  
- accuracy: 0.9829 - val\_loss: 0.0814 - val\_accuracy: 0.9760

Epoch 5/10

1875/1875 [=====] - 1s 659us/step - loss: 0.0454  
- accuracy: 0.9858 - val\_loss: 0.0796 - val\_accuracy: 0.9758

Epoch 6/10

1875/1875 [=====] - 1s 659us/step - loss: 0.0356  
- accuracy: 0.9887 - val\_loss: 0.0813 - val\_accuracy: 0.9760

Epoch 7/10

1875/1875 [=====] - 1s 657us/step - loss: 0.0287  
- accuracy: 0.9913 - val\_loss: 0.0810 - val\_accuracy: 0.9756

Epoch 8/10

1875/1875 [=====] - 1s 725us/step - loss: 0.0229  
- accuracy: 0.9928 - val\_loss: 0.0782 - val\_accuracy: 0.9775

Epoch 9/10

1875/1875 [=====] - 1s 673us/step - loss: 0.0182  
- accuracy: 0.9945 - val\_loss: 0.0813 - val\_accuracy: 0.9785

Epoch 10/10

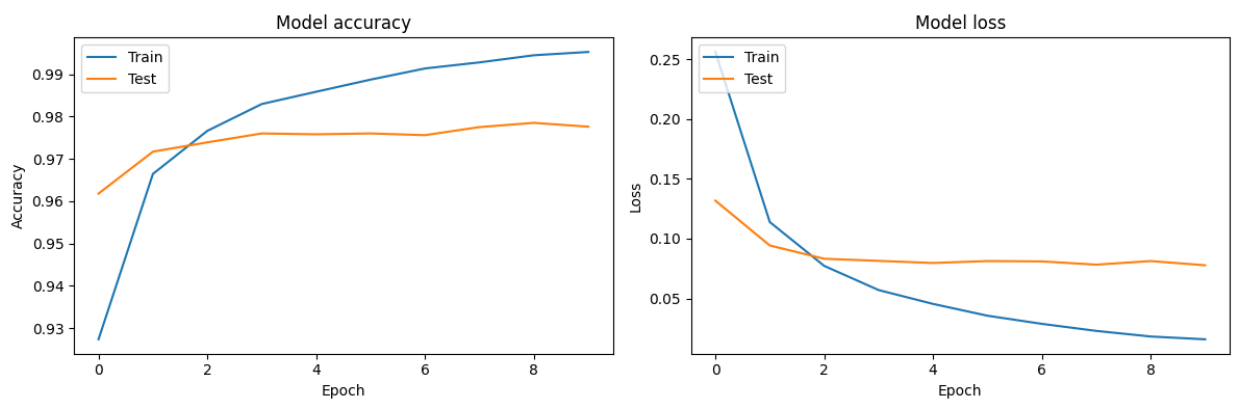
1875/1875 [=====] - 1s 657us/step - loss: 0.0158  
- accuracy: 0.9952 - val\_loss: 0.0777 - val\_accuracy: 0.9776

```
In [29]: # Plot training & validation accuracy values
plt.figure(figsize=(12, 4))

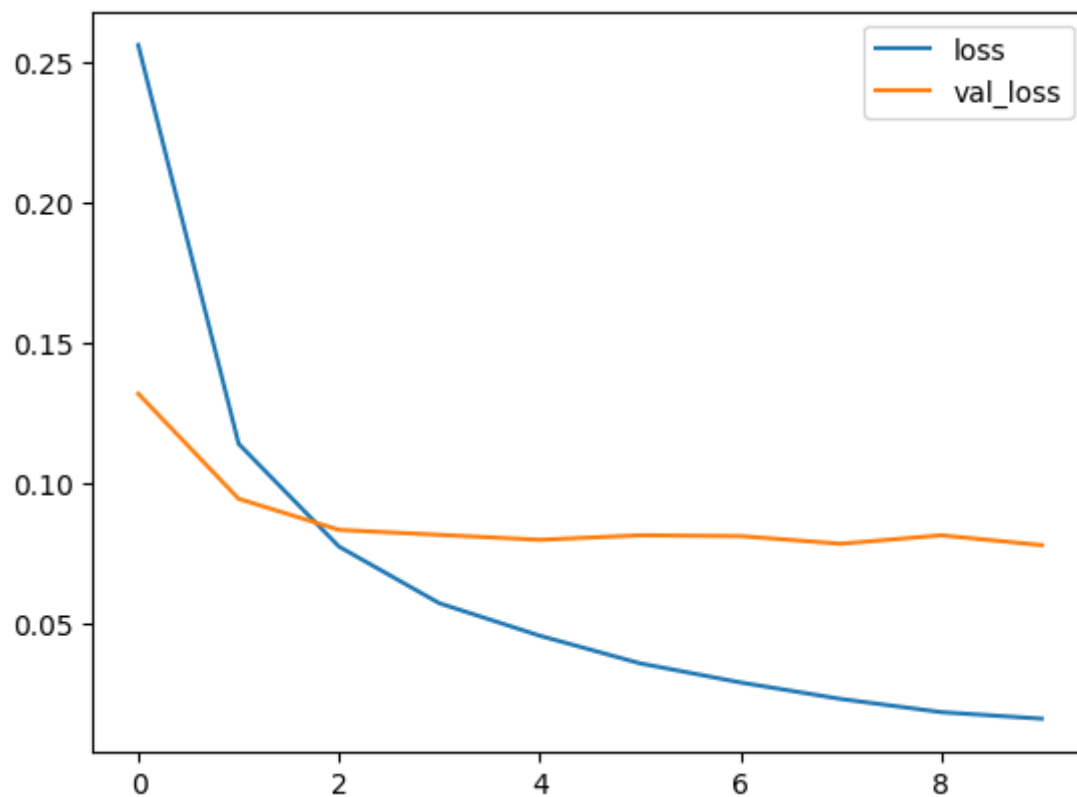
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')

# Plot training & validation loss values
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')

plt.tight_layout()
plt.show()
```



```
In [34]: history_df = pd.DataFrame(history.history)
history_df.loc[:, ['loss', 'val_loss']].plot();
```



```
In [19]: history.history
```

```
Out[19]: {'loss': [0.22070597112178802,
0.09154766798019409,
0.06130516901612282,
0.04478169232606888,
0.03293851763010025,
0.025184987112879753,
0.019530057907104492,
0.01601477712392807,
0.01281061489135027,
0.010680190287530422,
0.009442095644772053,
0.008974487893283367,
0.007534516975283623,
0.007101559080183506,
0.006281865295022726],
'accuracy': [0.9360666871070862,
0.9730833172798157,
0.9814833402633667,
0.9863499999046326,
0.989983320236206,
0.9916666746139526,
0.9935666918754578,
0.9950666427612305,
0.9958999752998352,
0.996483325958252,
0.9971833229064941,
0.9970333576202393,
0.9974666833877563,
0.9975166916847229,
0.9977666735649109],
'val_loss': [0.11028121411800385,
0.08852854371070862,
0.07909859716892242,
0.0747181624174118,
0.07619019597768784,
0.06984405219554901,
0.06666118651628494,
0.07518196105957031,
0.08322479575872421,
0.0868171975016594,
0.08774709701538086,
0.08541646599769592,
0.11101407557725906,
0.09484095126390457,
0.09372469782829285],
'val_accuracy': [0.9681000113487244,
0.9710999727249146,
0.9753999710083008,
0.975600004196167,
0.9771999716758728,
0.9801999926567078,
0.9805999994277954,
0.9793999791145325,
0.9779999852180481,
0.9785000085830688,
0.9782999753952026,
0.9807000160217285,
```

```
0.9757999777793884,  
0.9789000153541565,  
0.9801999926567078]}
```

well, it seems that loss and the accuracy still have some space to improve since they are not following to the trend of train.

## OK, now its time to see if this model acctually works,

### Lets Predict

Make Predictions:

Use the predict method of your trained model to get predictions. This method will return the predicted probabilities for each class

predictions will be an array with shape (number\_of\_samples, number\_of\_classes). Each row will contain the predicted probabilities for each class.

```
In [35]: predictions = model.predict(test_images)
```

```
313/313 [=====] - 0s 323us/step
```

Get Predicted Classes:

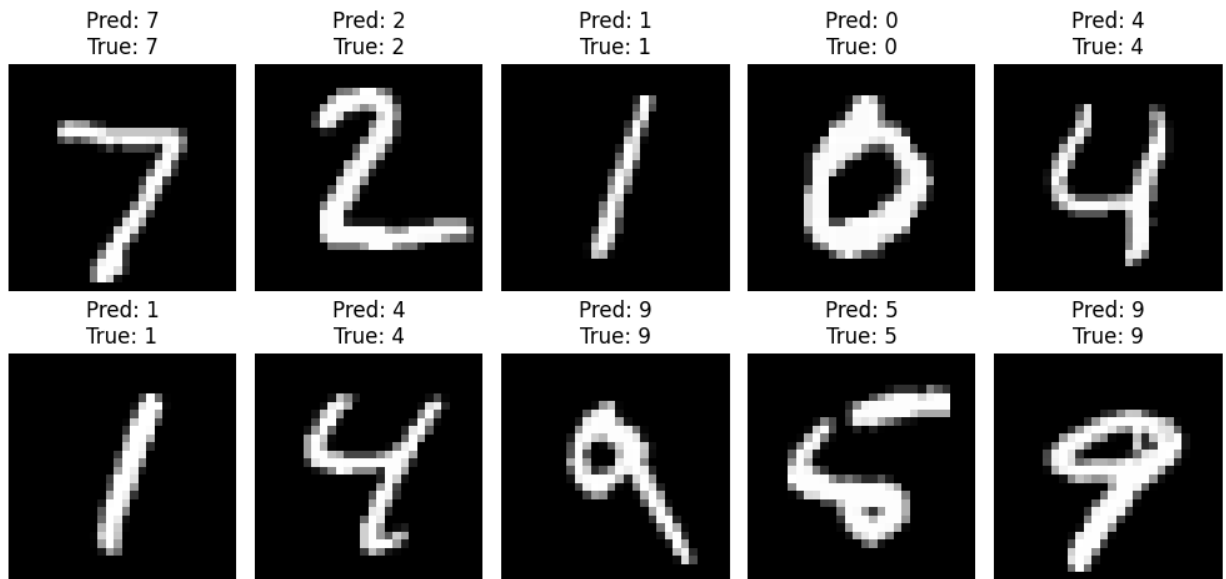
If we want to get the class with the highest probability for each sample (i.e., the predicted class label), we can use np.argmax:

```
In [36]: predicted_labels = np.argmax(predictions, axis=1)
```



## Visualize Predictions

```
In [37]: # Visualize the first 10 predictions
plt.figure(figsize=(10, 5))
for i in range(10):
    plt.subplot(2, 5, i+1)
    plt.imshow(test_images[i], cmap='gray')
    plt.title(f"Pred: {predicted_labels[i]}\nTrue: {np.argmax(test_labels[i])}")
    plt.axis('off')
plt.tight_layout()
plt.show()
```



## Evaluate Prediction Accuracy

```
In [38]: accuracy = np.mean(predicted_labels == np.argmax(test_labels, axis=1))
print(f"Prediction accuracy: {accuracy * 100:.2f}%")
```

Prediction accuracy: 97.76%

In [ ]: